

08\_16\_23

August 17, 2023

## 1 Imports/Device Settings/Utils

### 1.1 RUN ONLY ONCE FROM HERE

```
[59]: import numpy as np
import torch
import torch.nn as nn
#import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import utils

from PIL import Image
#from skimage.metrics import peak_signal_noise_ratio as psnr

device = (
    "cuda"
    if torch.cuda.is_available()
    else "cpu"
)

print(f"Using {device} device")
torch.set_default_device(device)
torch.set_default_dtype(torch.float32)

IMAGE_DIM = 48
```

Using cuda device

```
[60]: ground_truth_images = np.load("training.npy")
```

```
[61]: #ground_truth_images = np.zeros((28709, 1, IMAGE_DIM, IMAGE_DIM), np.float32)
#utils.create_training_data("train/", ground_truth_images)

#normalizations step
#ground_truth_images = np.multiply(ground_truth_images, 1/255)
#np.save("training.npy", ground_truth_images)
```

```
[62]: #revert normalization
ground_truth_images = np.multiply(ground_truth_images, 255)
clean = torch.as_tensor(ground_truth_images, dtype=torch.float32)
```

## 1.2 UNTIL HERE!!

### 1.3 Then back to relevant code:

```
[91]: # Create A.T*y
# CHOOSE SAMPLING RATE HERE

p = IMAGE_DIM**2
n = int(0.5*p)
A = np.random.normal(loc=0, scale=1/n, size=(n, p))

added_noise_images = np.zeros((28709, 1, IMAGE_DIM, IMAGE_DIM), np.float32)

i = 0
for image in ground_truth_images:
    if (i%1000==0):
        print(f"Image {i}")
        x = ground_truth_images[i][0]
        x = np.reshape(x, (IMAGE_DIM**2), 'F')

        added_noise_images[i][0] = np.reshape(np.matmul(A.T, np.matmul(A, x)),
        ↪(IMAGE_DIM, IMAGE_DIM), 'F')
        i += 1

print("done")
```

```
Image 0
Image 1000
Image 2000
Image 3000
Image 4000
Image 5000
Image 6000
Image 7000
Image 8000
Image 9000
Image 10000
Image 11000
Image 12000
Image 13000
Image 14000
Image 15000
Image 16000
Image 17000
```

```
Image 18000
Image 19000
Image 20000
Image 21000
Image 22000
Image 23000
Image 24000
Image 25000
Image 26000
Image 27000
Image 28000
done
```

```
[92]: noisy = torch.as_tensor(added_noise_images, dtype=torch.float32)

# Set training dictionary

train = {"noisy" : noisy[ : 25000],
        "clean" : clean[ : 25000]}

# Set validation dictionary
validation = {"noisy" : noisy[25000 : ],
             "clean" : clean[25000 : ]}
```

#### 1.4 relevant calculations:

```
[93]: # A priori PSNR calculation
test_number = len(train["noisy"])
print(f"# of test images: {test_number}")

ante_psnr = utils.avg_psnr(test_number, train["noisy"], train["clean"])
print(f"Avg. PSNR: {round(ante_psnr, 2)}")
```

```
# of test images: 25000
```



```
/burg/opt/anaconda3-2022.05/lib/python3.9/site-
packages/skimage/metrics/simple_metrics.py:163: RuntimeWarning: invalid value
encountered in double_scalars
  return 10 * np.log10((data_range ** 2) / err)
```

```
Avg. PSNR: 4.35
```

```
[94]: # Optional - compare noisy vs. clean image
sample_no = 1
im1 = utils.detach(noisy[sample_no][0])
im2 = utils.detach(clean[sample_no][0])
utils.show_images(im1, im2)

# Optional - test pixel distributions
```

```
utils.show_images_hist(im1, im2)
```



## 2 Correctness of image flattening

```
[95]: # temp = ground_truth_images[0][0]
      # column1 = temp[0:10, 0]
      # print(column1)

      # columns = np.reshape(temp, (2304), 'F')
      # print(columns[0:10])

[96]: # p = len(temp)
      # n = int(p/2)
      # print(p, n)

      # A = np.random.normal(loc=0, scale=1/n, size=(n, p))
      # y = np.matmul(A, temp)
      # Ay = np.matmul(A.T, y)
      # Ay = np.reshape(Ay, (48, 48), 'F')

      # true_image = ground_truth_images[0][0]
      # test_image = Ay
      # d_range = np.max([np.max(test_image) - np.min(test_image), np.max(true_image)
      ↪- np.min(true_image)])
      # print(psnr(true_image, test_image, data_range=d_range))
```

## 3 U-net architecture

```
[97]: # Establish neural network model

class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(ConvBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
        ↪stride=1, padding=1)
        self.actv1 = nn.ReLU()

    def forward(self, x):
        x = self.actv1(self.conv1(x))
        return x

# Establish neural network model

class Twonet(nn.Module):
    def __init__(self):
        super(Twonet, self).__init__()

        #1 if batched, 0 if unbatched
```

```

self.cat_dimension = 1

self.conv1_1 = ConvBlock(in_channels=1, out_channels=64)
self.conv1_2 = ConvBlock(in_channels=64, out_channels=64)

self.maxPool1 = nn.MaxPool2d(kernel_size=2)

self.conv2_1 = ConvBlock(in_channels=64, out_channels=128)
self.conv2_2 = ConvBlock(in_channels=128, out_channels=128)

self.maxPool2 = nn.MaxPool2d(kernel_size=2)

self.conv3_1 = ConvBlock(in_channels=128, out_channels=256)
self.conv3_2 = ConvBlock(in_channels=256, out_channels=256)

self.convTranspose1 = nn.ConvTranspose2d(in_channels=256,
↪out_channels=128, kernel_size=2, stride=2)

self.conv4_1 = ConvBlock(in_channels=256, out_channels=128)
self.conv4_2 = ConvBlock(in_channels=128, out_channels=128)

self.convTranspose2 = nn.ConvTranspose2d(in_channels=128,
↪out_channels=64, kernel_size=2, stride=2)

self.conv5_1 = ConvBlock(in_channels=128, out_channels=64)
self.conv5_2 = ConvBlock(in_channels=64, out_channels=1)

#DONT FORGET TO ADD CONCATENATION CHANNELS

def forward(self, x):
    x = self.conv1_2(self.conv1_1(x))

    skip_connect1 = x

    x = self.maxPool1(x)
    x = self.conv2_2(self.conv2_1(x))

    skip_connect2 = x

    x = self.maxPool2(x)
    x = self.conv3_2(self.conv3_1(x))

    x = self.convTranspose1(x)

    x = torch.cat([skip_connect2, x], self.cat_dimension)
    x = self.conv4_2(self.conv4_1(x))

```

```

        x = self.convTranspose2(x)

        x = torch.cat([skip_connect1, x], self.cat_dimension)
        x = self.conv5_2(self.conv5_1(x))

        return x

network = Twonet()
network = network.to(device)
network.train()

```

```

[97]: Twonet(
  (conv1_1): ConvBlock(
    (conv1): Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (conv1_2): ConvBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (maxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (conv2_1): ConvBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (conv2_2): ConvBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (maxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (conv3_1): ConvBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (conv3_2): ConvBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (convTranspose1): ConvTranspose2d(256, 128, kernel_size=(2, 2), stride=(2, 2))
  (conv4_1): ConvBlock(
    (conv1): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (conv4_2): ConvBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

        (actv1): ReLU()
    )
    (convTranspose2): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2))
    (conv5_1): ConvBlock(
        (conv1): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (conv5_2): ConvBlock(
        (conv1): Conv2d(64, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
)

```

## 4 Optional Tests

```

[98]: # OPTIONAL
      # Testing input/output sizes
      input = torch.randn(2, 1, IMAGE_DIM, IMAGE_DIM)
      print("Input: ", np.shape(input))

      output1 = network(input)
      print("UNet Output: ", np.shape(output1))

```

```

Input:  torch.Size([2, 1, 48, 48])
UNet Output:  torch.Size([2, 1, 48, 48])

```

## 5 Training the Model

### 5.0.1 Set hyperparameters

```

[99]: epochs = 50
      batch_size = 20
      lr = 1e-4
      loss_fn = nn.MSELoss()
      optimizer = torch.optim.Adam(network.parameters(), lr=lr)

```

```

[100]: # Training
       training_loss, test_loss = utils.train(
           x=train["noisy"],
           y=train["clean"],
           validation_x=validation["noisy"],
           validation_y=validation["clean"],
           neural_network=network,
           epochs=epochs,
           batch_size=batch_size,
           learning_rate=lr,

```



```
    loss_function=loss_fn,  
    optimizer=optimizer  
)
```

Starting training...


```
epoch 0 complete. elapsed time: 7s  
training loss: 1861.3309326171875, test loss: 735.1430053710938  
epoch 1 complete. elapsed time: 14s  
training loss: 730.4447631835938, test loss: 671.6201782226562  
epoch 2 complete. elapsed time: 21s  
training loss: 659.782470703125, test loss: 612.9658203125  
epoch 3 complete. elapsed time: 28s  
training loss: 613.9469604492188, test loss: 565.3809814453125  
epoch 4 complete. elapsed time: 35s  
training loss: 578.6698608398438, test loss: 536.9165649414062  
epoch 5 complete. elapsed time: 42s  
training loss: 549.6648559570312, test loss: 511.2002868652344  
epoch 6 complete. elapsed time: 49s  
training loss: 526.0184326171875, test loss: 483.69854736328125  
epoch 7 complete. elapsed time: 56s  
training loss: 504.69744873046875, test loss: 465.8990173339844  
epoch 8 complete. elapsed time: 63s  
training loss: 483.731689453125, test loss: 446.2145080566406  
epoch 9 complete. elapsed time: 70s  
training loss: 463.05633544921875, test loss: 431.4891662597656  
epoch 10 complete. elapsed time: 77s  
training loss: 444.573974609375, test loss: 417.3224182128906  
epoch 11 complete. elapsed time: 84s  
training loss: 428.67437744140625, test loss: 403.796630859375  
epoch 12 complete. elapsed time: 91s  
training loss: 414.38751220703125, test loss: 394.0511474609375  
epoch 13 complete. elapsed time: 98s  
training loss: 401.9185791015625, test loss: 383.3927307128906  
epoch 14 complete. elapsed time: 105s  
training loss: 390.56622314453125, test loss: 374.6682434082031  
epoch 15 complete. elapsed time: 112s  
training loss: 380.46270751953125, test loss: 366.3709716796875  
epoch 16 complete. elapsed time: 119s  
training loss: 371.0546875, test loss: 360.0040283203125  
epoch 17 complete. elapsed time: 126s  
training loss: 362.37994384765625, test loss: 352.8594665527344  
epoch 18 complete. elapsed time: 133s  
training loss: 354.3267517089844, test loss: 346.6192321777344  
epoch 19 complete. elapsed time: 140s  
training loss: 346.608154296875, test loss: 340.80279541015625  
epoch 20 complete. elapsed time: 147s  
training loss: 339.265380859375, test loss: 334.51715087890625  
epoch 21 complete. elapsed time: 154s
```

training loss: 332.2876281738281, test loss: 327.9779052734375  
epoch 22 complete. elapsed time: 162s  
training loss: 325.6748352050781, test loss: 322.1311340332031  
epoch 23 complete. elapsed time: 169s  
training loss: 319.4094543457031, test loss: 316.9410705566406  
epoch 24 complete. elapsed time: 176s  
training loss: 313.5443420410156, test loss: 311.3306884765625  
epoch 25 complete. elapsed time: 183s  
training loss: 308.10992431640625, test loss: 305.7515563964844  
epoch 26 complete. elapsed time: 190s  
training loss: 302.8777160644531, test loss: 300.5751037597656  
epoch 27 complete. elapsed time: 197s  
training loss: 297.8593444824219, test loss: 296.4984130859375  
epoch 28 complete. elapsed time: 204s  
training loss: 292.9192199707031, test loss: 292.3488464355469  
epoch 29 complete. elapsed time: 211s  
training loss: 288.485595703125, test loss: 288.4309997558594  
epoch 30 complete. elapsed time: 219s  
training loss: 284.235107421875, test loss: 285.108642578125  
epoch 31 complete. elapsed time: 226s  
training loss: 280.0912780761719, test loss: 281.8576965332031  
epoch 32 complete. elapsed time: 234s  
training loss: 276.2082824707031, test loss: 278.7071533203125  
epoch 33 complete. elapsed time: 241s  
training loss: 272.52239990234375, test loss: 275.643798828125  
epoch 34 complete. elapsed time: 248s  
training loss: 269.068603515625, test loss: 272.75372314453125  
epoch 35 complete. elapsed time: 254s  
training loss: 265.6302185058594, test loss: 269.83203125  
epoch 36 complete. elapsed time: 261s  
training loss: 262.4419860839844, test loss: 267.4181823730469  
epoch 37 complete. elapsed time: 268s  
training loss: 259.359130859375, test loss: 264.8430480957031  
epoch 38 complete. elapsed time: 275s  
training loss: 256.4273376464844, test loss: 262.6467590332031  
epoch 39 complete. elapsed time: 282s  
training loss: 253.5440673828125, test loss: 260.4901428222656  
epoch 40 complete. elapsed time: 289s  
training loss: 250.83164978027344, test loss: 258.5862731933594  
epoch 41 complete. elapsed time: 296s  
training loss: 248.2327880859375, test loss: 256.67889404296875  
epoch 42 complete. elapsed time: 303s  
training loss: 245.7191619873047, test loss: 254.65151977539062  
epoch 43 complete. elapsed time: 310s  
training loss: 243.3513641357422, test loss: 252.72213745117188  
epoch 44 complete. elapsed time: 317s  
training loss: 240.98834228515625, test loss: 251.26190185546875  
epoch 45 complete. elapsed time: 324s

```
training loss: 238.8385467529297, test loss: 249.60926818847656
epoch 46 complete. elapsed time: 330s
training loss: 236.67416381835938, test loss: 247.65902709960938
epoch 47 complete. elapsed time: 337s
training loss: 234.61024475097656, test loss: 245.51751708984375
epoch 48 complete. elapsed time: 344s
training loss: 232.65182495117188, test loss: 243.7272186279297
epoch 49 complete. elapsed time: 351s
training loss: 230.67630004882812, test loss: 242.05047607421875
Done training.
```

## 6 Loss graph

```
[101]: epoch_record = np.arange(0, len(training_loss), 1)
plt.plot(epoch_record, training_loss, c="red", label="training MSE")
plt.plot(epoch_record, test_loss, c="blue", label="test MSE")
plt.legend(loc="upper right")
plt.title(f"Training loss over {epochs} epochs")
plt.xlabel("Epochs")
plt.ylabel("MSE")
plt.show()
```



## 7 Tests/PSNR Calculations

```
[102]: network.eval()
```

```
[102]: Twonet(  
    (conv1_1): ConvBlock(  
        (conv1): Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (conv1_2): ConvBlock(  
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (maxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (conv2_1): ConvBlock(  
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (conv2_2): ConvBlock(  
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (maxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceil_mode=False)  
    (conv3_1): ConvBlock(  
        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (conv3_2): ConvBlock(  
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (convTranspose1): ConvTranspose2d(256, 128, kernel_size=(2, 2), stride=(2, 2))  
    (conv4_1): ConvBlock(  
        (conv1): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (conv4_2): ConvBlock(  
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
    (convTranspose2): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2))  
    (conv5_1): ConvBlock(  
        (conv1): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (actv1): ReLU()  
    )  
)
```

```

(conv5_2): ConvBlock(
  (conv1): Conv2d(64, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
)

```

```

[103]: # Choose between [0, 19]
sample_number = 1
test_size = 100

# Test visual recovery results
with torch.no_grad():
    predictions = network(validation["noisy"][0:test_size])

post_psnr = utils.avg_psnr(test_size, predictions, validation["clean"])
print(f"Avg. PSNR: {round(post_psnr, 2)}")


```

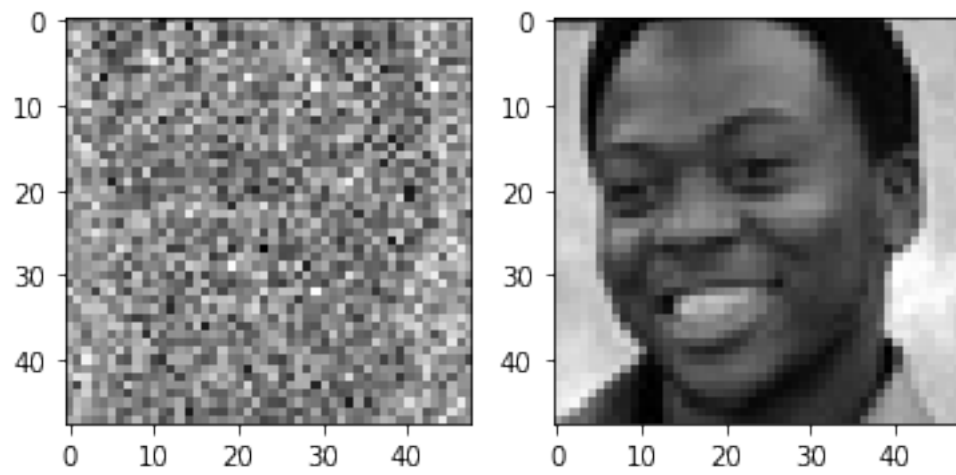
Avg. PSNR: 24.08

```

[104]: sample_no = 29
im0 = utils.detach(validation["noisy"][sample_no][0])
im1 = utils.detach(predictions[sample_no][0])
im2 = utils.detach(validation["clean"][sample_no][0])
utils.show_images(im0, im1)
utils.show_images(im0, im2)

```





[ ]: