

08\_16\_23

August 17, 2023

# 1 Imports/Device Settings/Utils

## 1.1 RUN ONLY ONCE FROM HERE

```
[1]: import numpy as np
import torch
import torch.nn as nn
#import torchvision.transforms as transforms
import matplotlib.pyplot as plt
import utils

from PIL import Image
#from skimage.metrics import peak_signal_noise_ratio as psnr

device = (
    "cuda"
    if torch.cuda.is_available()
    else "cpu"
)

print(f"Using {device} device")
torch.set_default_device(device)
torch.set_default_dtype(torch.float32)

IMAGE_DIM = 48
```

Using cuda device

```
[2]: ground_truth_images = np.load("training.npy")
```

```
[3]: #ground_truth_images = np.zeros((28709, 1, IMAGE_DIM, IMAGE_DIM), np.float32)
#utils.create_training_data("train/", ground_truth_images)

#normalizations step
#ground_truth_images = np.multiply(ground_truth_images, 1/255)
#np.save("training.npy", ground_truth_images)
```

```
[4]: #revert normalization
ground_truth_images = np.multiply(ground_truth_images, 255)
clean = torch.as_tensor(ground_truth_images, dtype=torch.float32)
```

## 1.2 UNTIL HERE!!

### 1.3 Then back to relevant code:

```
[5]: # Create A.T*y
# CHOOSE SAMPLING RATE HERE

p = IMAGE_DIM**2
n = int(0.5*p)
A = np.random.normal(loc=0, scale=1/n, size=(n, p))

added_noise_images = np.zeros((28709, 1, IMAGE_DIM, IMAGE_DIM), np.float32)

i = 0
for image in ground_truth_images:
    if (i%1000==0):
        print(f"Image {i}")
        x = ground_truth_images[i][0]
        x = np.reshape(x, (IMAGE_DIM**2), 'F')

        added_noise_images[i][0] = np.reshape(np.matmul(A.T, np.matmul(A, x)),
        ↪(IMAGE_DIM, IMAGE_DIM), 'F')
        i += 1

print("done")
```

```
Image 0
Image 1000
Image 2000
Image 3000
Image 4000
Image 5000
Image 6000
Image 7000
Image 8000
Image 9000
Image 10000
Image 11000
Image 12000
Image 13000
Image 14000
Image 15000
Image 16000
Image 17000
```

```

Image 18000
Image 19000
Image 20000
Image 21000
Image 22000
Image 23000
Image 24000
Image 25000
Image 26000
Image 27000
Image 28000
done

```

```

[6]: noisy = torch.as_tensor(added_noise_images, dtype=torch.float32)

# Set training dictionary

train = {"noisy" : noisy[ : 25000],
        "clean" : clean[ : 25000]}

# Set validation dictionary
validation = {"noisy" : noisy[25000 : ],
             "clean" : clean[25000 : ]}

```

#### 1.4 relevant calculations:

```

[7]: # A priori PSNR calculation
test_number = len(train["noisy"])
print(f"# of test images: {test_number}")

ante_psnr = utils.avg_psnr(test_number, train["noisy"], train["clean"])
print(f"Avg. PSNR: {round(ante_psnr, 2)}")

# of test images: 25000

/burg/opt/anaconda3-2022.05/lib/python3.9/site-
packages/skimage/metrics/simple_metrics.py:163: RuntimeWarning: invalid value
encountered in double_scalars
  return 10 * np.log10((data_range ** 2) / err)

Avg. PSNR: 4.35

```

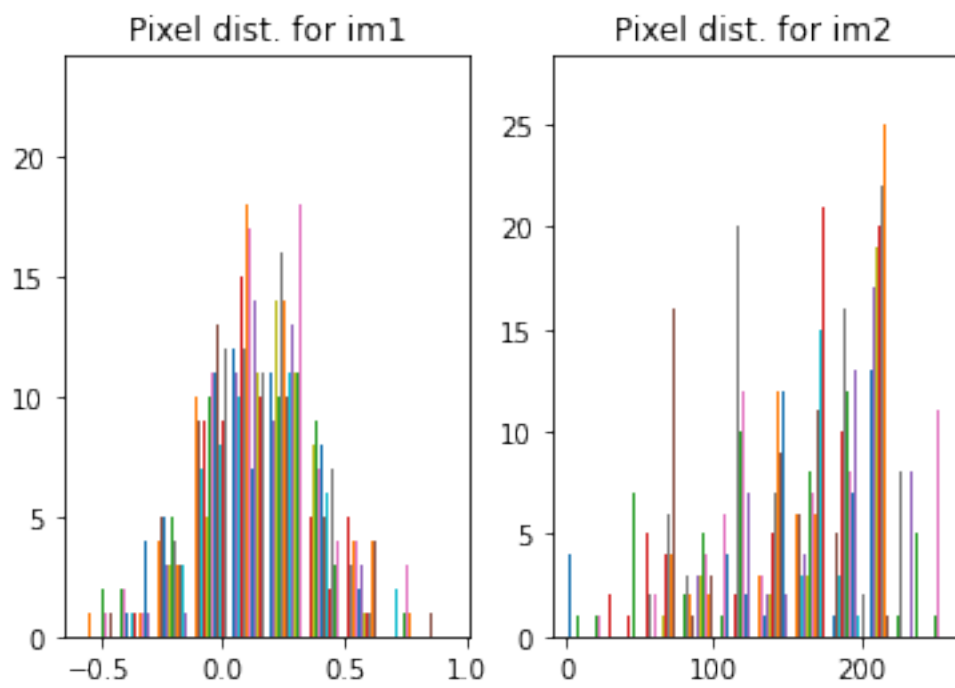
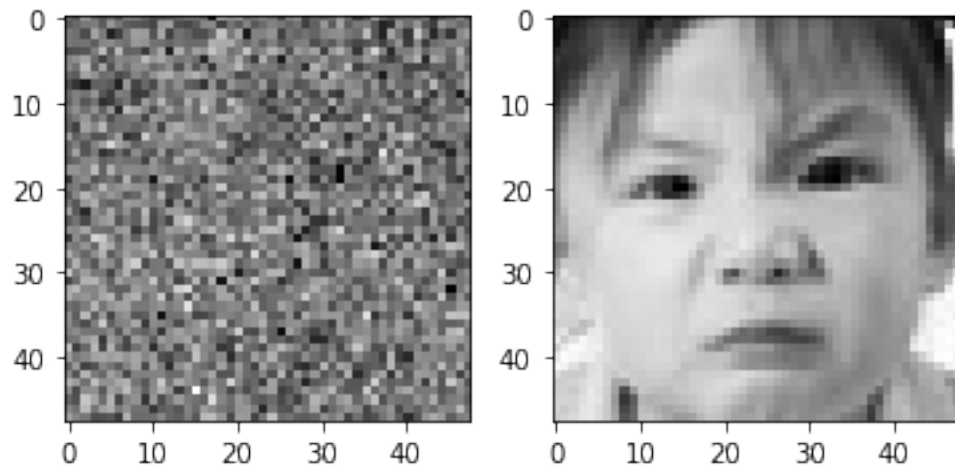
```

[8]: # Optional - compare noisy vs. clean image
sample_no = 1
im1 = utils.detach(noisy[sample_no][0])
im2 = utils.detach(clean[sample_no][0])
utils.show_images(im1, im2)

# Optional - test pixel distributions

```

```
utils.show_images_hist(im1, im2)
```



## 2 Correctness of image flattening

```
[9]: # temp = ground_truth_images[0][0]
# column1 = temp[0:10, 0]
# print(column1)

# columns = np.reshape(temp, (2304), 'F')
# print(columns[0:10])
```

```
[10]: # p = len(temp)
# n = int(p/2)
# print(p, n)

# A = np.random.normal(loc=0, scale=1/n, size=(n, p))
# y = np.matmul(A, temp)
# Ay = np.matmul(A.T, y)
# Ay = np.reshape(Ay, (48, 48), 'F')

# true_image = ground_truth_images[0][0]
# test_image = Ay
# d_range = np.max([np.max(test_image) - np.min(test_image), np.max(true_image)
# ↪- np.min(true_image)])
# print(psnr(true_image, test_image, data_range=d_range))
```

## 3 U-net architecture

```
[11]: # Establish neural network model

class ConvBlock(nn.Module):
    def __init__(self, in_channels, out_channels):
        super(ConvBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
↪stride=1, padding=1)
        self.actv1 = nn.ReLU()

    def forward(self, x):
        x = self.actv1(self.conv1(x))
        return x

# Establish neural network model

class Threenet(nn.Module):
    def __init__(self):
        super(Threenet, self).__init__()

        #1 if batched, 0 if unbatched
```

```

self.cat_dimension = 1

self.conv1_1 = ConvBlock(in_channels=1, out_channels=64)
self.conv1_2 = ConvBlock(in_channels=64, out_channels=64)

self.maxPool1 = nn.MaxPool2d(kernel_size=2)

self.conv2_1 = ConvBlock(in_channels=64, out_channels=128)
self.conv2_2 = ConvBlock(in_channels=128, out_channels=128)

self.maxPool2 = nn.MaxPool2d(kernel_size=2)

self.conv3_1 = ConvBlock(in_channels=128, out_channels=256)
self.conv3_2 = ConvBlock(in_channels=256, out_channels=256)

self.maxPool3 = nn.MaxPool2d(kernel_size=2)

self.conv4_1 = ConvBlock(in_channels=256, out_channels=512)
self.conv4_2 = ConvBlock(in_channels=512, out_channels=512)

self.convTranspose1 = nn.ConvTranspose2d(in_channels=512,
↪out_channels=256, kernel_size=2, stride=2)

self.conv5_1 = ConvBlock(in_channels=512, out_channels=256)
self.conv5_2 = ConvBlock(in_channels=256, out_channels=256)

self.convTranspose2 = nn.ConvTranspose2d(in_channels=256,
↪out_channels=128, kernel_size=2, stride=2)

self.conv6_1 = ConvBlock(in_channels=256, out_channels=128)
self.conv6_2 = ConvBlock(in_channels=128, out_channels=128)

self.convTranspose3 = nn.ConvTranspose2d(in_channels=128,
↪out_channels=64, kernel_size=2, stride=2)

self.conv7_1 = ConvBlock(in_channels=128, out_channels=64)
self.conv7_2 = ConvBlock(in_channels=64, out_channels=64)
self.conv7_3 = ConvBlock(in_channels=64, out_channels=1)

#DONT FORGET TO ADD CONCATENATION CHANNELS

def forward(self, x):
    x = self.conv1_2(self.conv1_1(x))

    skip_connect1 = x

    x = self.maxPool1(x)

```

```

        x = self.conv2_2(self.conv2_1(x))

        skip_connect2 = x

        x = self.maxPool2(x)
        x = self.conv3_2(self.conv3_1(x))

        skip_connect3 = x

        x = self.maxPool3(x)
        x = self.conv4_2(self.conv4_1(x))

        x = self.convTranspose1(x)
        x = torch.cat([skip_connect3, x], self.cat_dimension)
        x = self.conv5_2(self.conv5_1(x))

        x = self.convTranspose2(x)
        x = torch.cat([skip_connect2, x], self.cat_dimension)
        x = self.conv6_2(self.conv6_1(x))

        x = self.convTranspose3(x)
        x = torch.cat([skip_connect1, x], self.cat_dimension)
        x = self.conv7_3(self.conv7_2(self.conv7_1(x)))

    return x

network = Threenet()
network = network.to(device)
network.train()

```

```

[11]: Threenet(
  (conv1_1): ConvBlock(
    (conv1): Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (conv1_2): ConvBlock(
    (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (maxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (conv2_1): ConvBlock(
    (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (actv1): ReLU()
  )
  (conv2_2): ConvBlock(
    (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

    (actv1): ReLU()
)
(maxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(conv3_1): ConvBlock(
  (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv3_2): ConvBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(maxPool3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(conv4_1): ConvBlock(
  (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv4_2): ConvBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(convTranspose1): ConvTranspose2d(512, 256, kernel_size=(2, 2), stride=(2, 2))
(conv5_1): ConvBlock(
  (conv1): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv5_2): ConvBlock(
  (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(convTranspose2): ConvTranspose2d(256, 128, kernel_size=(2, 2), stride=(2, 2))
(conv6_1): ConvBlock(
  (conv1): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv6_2): ConvBlock(
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(convTranspose3): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2))
(conv7_1): ConvBlock(
  (conv1): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv7_2): ConvBlock(
  (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```



```

        (activ1): ReLU()
    )
    (conv7_3): ConvBlock(
      (conv1): Conv2d(64, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
      (activ1): ReLU()
    )
  )
)

```

## 4 Optional Tests

```

[12]: # OPTIONAL
      # Testing input/output sizes
      input = torch.randn(2, 1, IMAGE_DIM, IMAGE_DIM)
      print("Input: ", np.shape(input))

      output1 = network(input)
      print("UNet Output: ", np.shape(output1))

```

```

Input:  torch.Size([2, 1, 48, 48])
UNet Output:  torch.Size([2, 1, 48, 48])

```

## 5 Training the Model

### 5.0.1 Set hyperparameters

```

[13]: epochs = 50
      batch_size = 20
      lr = 1e-4
      loss_fn = nn.MSELoss()
      optimizer = torch.optim.Adam(network.parameters(), lr=lr)

```

```

[14]: # Training
      training_loss, test_loss = utils.train(
          x=train["noisy"],
          y=train["clean"],
          validation_x=validation["noisy"],
          validation_y=validation["clean"],
          neural_network=network,
          epochs=epochs,
          batch_size=batch_size,
          learning_rate=lr,
          loss_function=loss_fn,
          optimizer=optimizer
      )

```

```

Starting training...
epoch 0 complete. elapsed time: 15s

```


training loss: 1802.5513916015625, test loss: 665.0151977539062  
epoch 1 complete. elapsed time: 25s  
training loss: 622.411865234375, test loss: 523.985107421875  
epoch 2 complete. elapsed time: 35s  
training loss: 516.7144165039062, test loss: 465.3548583984375  
epoch 3 complete. elapsed time: 45s  
training loss: 461.79217529296875, test loss: 422.3830871582031  
epoch 4 complete. elapsed time: 55s  
training loss: 422.38433837890625, test loss: 393.26788330078125  
epoch 5 complete. elapsed time: 65s  
training loss: 395.4277038574219, test loss: 371.1346130371094  
epoch 6 complete. elapsed time: 75s  
training loss: 374.3458557128906, test loss: 350.6646423339844  
epoch 7 complete. elapsed time: 84s  
training loss: 357.1914367675781, test loss: 337.5110778808594  
epoch 8 complete. elapsed time: 94s  
training loss: 341.9107666015625, test loss: 324.4728088378906  
epoch 9 complete. elapsed time: 104s  
training loss: 328.2871398925781, test loss: 313.4806823730469  
epoch 10 complete. elapsed time: 114s  
training loss: 316.0379333496094, test loss: 303.58551025390625  
epoch 11 complete. elapsed time: 124s  
training loss: 304.9957275390625, test loss: 294.5194396972656  
epoch 12 complete. elapsed time: 134s  
training loss: 295.2705078125, test loss: 286.2991943359375  
epoch 13 complete. elapsed time: 143s  
training loss: 286.2917785644531, test loss: 279.5648498535156  
epoch 14 complete. elapsed time: 153s  
training loss: 278.4056396484375, test loss: 273.3607177734375  
epoch 15 complete. elapsed time: 163s  
training loss: 271.0382385253906, test loss: 268.5890808105469  
epoch 16 complete. elapsed time: 173s  
training loss: 264.4971923828125, test loss: 263.87164306640625  
epoch 17 complete. elapsed time: 183s  
training loss: 258.1697692871094, test loss: 258.84442138671875  
epoch 18 complete. elapsed time: 193s  
training loss: 252.35044860839844, test loss: 255.31797790527344  
epoch 19 complete. elapsed time: 203s  
training loss: 246.97091674804688, test loss: 251.5403289794922  
epoch 20 complete. elapsed time: 213s  
training loss: 241.95431518554688, test loss: 249.09732055664062  
epoch 21 complete. elapsed time: 222s  
training loss: 237.2635955810547, test loss: 246.48399353027344  
epoch 22 complete. elapsed time: 232s  
training loss: 232.85414123535156, test loss: 243.36354064941406  
epoch 23 complete. elapsed time: 242s  
training loss: 228.5993194580078, test loss: 240.6340789794922  
epoch 24 complete. elapsed time: 252s

training loss: 224.567138671875, test loss: 237.85043334960938  
epoch 25 complete. elapsed time: 262s  
training loss: 220.77330017089844, test loss: 235.79922485351562  
epoch 26 complete. elapsed time: 272s  
training loss: 217.12429809570312, test loss: 233.82919311523438  
epoch 27 complete. elapsed time: 282s  
training loss: 213.58514404296875, test loss: 232.19166564941406  
epoch 28 complete. elapsed time: 292s  
training loss: 210.30291748046875, test loss: 229.68768310546875  
epoch 29 complete. elapsed time: 301s  
training loss: 207.23300170898438, test loss: 227.72613525390625  
epoch 30 complete. elapsed time: 311s  
training loss: 204.1021728515625, test loss: 225.81854248046875  
epoch 31 complete. elapsed time: 321s  
training loss: 201.107421875, test loss: 224.1342010498047  
epoch 32 complete. elapsed time: 331s  
training loss: 198.2167510986328, test loss: 222.49327087402344  
epoch 33 complete. elapsed time: 341s  
training loss: 195.44253540039062, test loss: 220.91770935058594  
epoch 34 complete. elapsed time: 350s  
training loss: 192.7393798828125, test loss: 219.41827392578125  
epoch 35 complete. elapsed time: 360s  
training loss: 190.14138793945312, test loss: 218.14453125  
epoch 36 complete. elapsed time: 370s  
training loss: 187.79193115234375, test loss: 216.89280700683594  
epoch 37 complete. elapsed time: 379s  
training loss: 185.7062225341797, test loss: 215.64926147460938  
epoch 38 complete. elapsed time: 389s  
training loss: 184.44827270507812, test loss: 214.45855712890625  
epoch 39 complete. elapsed time: 398s  
training loss: 183.015625, test loss: 213.6410675048828  
epoch 40 complete. elapsed time: 408s  
training loss: 180.52981567382812, test loss: 213.08062744140625  
epoch 41 complete. elapsed time: 417s  
training loss: 177.99095153808594, test loss: 212.54307556152344  
epoch 42 complete. elapsed time: 427s  
training loss: 175.25914001464844, test loss: 211.60459899902344  
epoch 43 complete. elapsed time: 437s  
training loss: 172.98565673828125, test loss: 211.47093200683594  
epoch 44 complete. elapsed time: 446s  
training loss: 170.97291564941406, test loss: 210.9196014404297  
epoch 45 complete. elapsed time: 456s  
training loss: 169.19891357421875, test loss: 211.03256225585938  
epoch 46 complete. elapsed time: 465s  
training loss: 167.78114318847656, test loss: 210.12930297851562  
epoch 47 complete. elapsed time: 475s  
training loss: 166.7620391845703, test loss: 210.3343963623047  
epoch 48 complete. elapsed time: 484s

```
training loss: 165.90383911132812, test loss: 210.46292114257812
epoch 49 complete. elapsed time: 494s
training loss: 165.0670166015625, test loss: 210.6776123046875
Done training.
```

## 6 Loss graph

```
[15]: epoch_record = np.arange(0, len(training_loss), 1)
plt.plot(epoch_record, training_loss, c="red", label="training MSE")
plt.plot(epoch_record, test_loss, c="blue", label="test MSE")
plt.legend(loc="upper right")
plt.title(f"Training loss over {epochs} epochs")
plt.xlabel("Epochs")
plt.ylabel("MSE")
plt.show()
```



## 7 Tests/PSNR Calculations

```
[16]: network.eval()
```

```
[16]: Threenet(
      (conv1_1): ConvBlock(
```

```

        (conv1): Conv2d(1, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (conv1_2): ConvBlock(
        (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (maxPool1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (conv2_1): ConvBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (conv2_2): ConvBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (maxPool2): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (conv3_1): ConvBlock(
        (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (conv3_2): ConvBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (maxPool3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (conv4_1): ConvBlock(
        (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (conv4_2): ConvBlock(
        (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (convTranspose1): ConvTranspose2d(512, 256, kernel_size=(2, 2), stride=(2, 2))
    (conv5_1): ConvBlock(
        (conv1): Conv2d(512, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (conv5_2): ConvBlock(
        (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
        (actv1): ReLU()
    )
    (convTranspose2): ConvTranspose2d(256, 128, kernel_size=(2, 2), stride=(2, 2))

```

```

(conv6_1): ConvBlock(
  (conv1): Conv2d(256, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv6_2): ConvBlock(
  (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(convTranspose3): ConvTranspose2d(128, 64, kernel_size=(2, 2), stride=(2, 2))
(conv7_1): ConvBlock(
  (conv1): Conv2d(128, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv7_2): ConvBlock(
  (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
(conv7_3): ConvBlock(
  (conv1): Conv2d(64, 1, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (actv1): ReLU()
)
)

```

```

[17]: # Choose between [0, 19]
sample_number = 1
test_size = 100

# Test visual recovery results
with torch.no_grad():
    predictions = network(validation["noisy"][0:test_size])

post_psnr = utils.avg_psnr(test_size, predictions, validation["clean"])
print(f"Avg. PSNR: {round(post_psnr, 2)}")



```

Avg. PSNR: 24.84

```

[18]: sample_no = 29
im0 = utils.detach(validation["noisy"][sample_no][0])
im1 = utils.detach(predictions[sample_no][0])
im2 = utils.detach(validation["clean"][sample_no][0])
utils.show_images(im0, im1)
utils.show_images(im0, im2)

```



[ ]: